

AI 534: Machine Learning HW1: Feature Map and k -NN (20%+3% extra)

Instructions:

1. In this HW you'll join a Kaggle competition hosted by us at <https://www.kaggle.com/t/775d0952c73240e1a6e7918bab76fe7e>. Use your OSU email for your account and **the last 5 digits of your OSU ID** as “team name”.
2. Suggested Work Plan and Participation in Kaggle:
 - by week 0: download data, register on Kaggle, and submit our `sample_submission.csv` to get onboard
 - by week 1: finish Parts 1-2 and Quiz 1; submit your best result to Kaggle (0.5 extra pts)
 - by week 2: finish Part 3 and Quiz 2; submit your new best result to Kaggle (0.5 extra); start on Part 4
 - by week 3: finish Parts 4-5; submit your overall best result to Kaggle; write report

3. This HW, like all other HWs, should be done in Python 3 and numpy. See the course homepage for a numpy tutorial (under Unit 1). If you don't have a Python+numpy installation yourself, you can use the College servers:

```
ssh username@access.engr.oregonstate.edu
```

replacing `username` with your ENGR user name (should be identical to your ONID login). See the following for instructions on using SSH from Windows:

<https://it.engineering.oregonstate.edu/accessing-unix-server-using-putty-ssh>

If you don't have an ENGR account, see <https://it.engineering.oregonstate.edu/get-engr-account>.

You're **highly recommended** to set up SSH keys to bypass Duo authentication and password:

<https://it.engineering.oregonstate.edu/ssh-keygen>

The default `python3` on ENGR servers have `numpy`, `matplotlib`, `sklearn`, and `pandas` installed.

4. Besides machine learning, this HW also teaches you data (pre-)processing skills and Unix (Linux or Mac OS X) command lines tools. These skills are even more important than machine learning itself for a software engineer or data scientist. As stated in the syllabus, Windows is **not** recommended. The instructor and the TA do not have access to Windows computers and will not be able to provide technical assistance for Windows. Use `cmd`, `cygwin`, or `ssh` (see above) if you don't have Linux or Mac OS X on your own computer.

For Linux, please read Ch. 1 of <https://buildmedia.readthedocs.org/media/pdf/lym/latest/lym.pdf>

If you prefer to use Jupyter notebook but have difficulty setting it up on your own computer, the TAs have written a tutorial on running jupyter notebook remotely on the server:

https://web.engr.oregonstate.edu/~huanlian/teaching/ML/2025fall/extra/TAs_jupyterTutorial.pdf

5. Ask questions on [Ed Discussion](#). You're encouraged to answer other students' questions as well.
6. Download HW1 data from either Canvas or the link below. Unzip it (terminals: `$ tar -xzvf hw1-data.tgz`; here `-x` means “extract” and `z` means “(un)zip”; `$` is the prompt; Windows: WinZip or 7-zip). It contains:

<code>income.train.5k.csv</code>	training data (5,000 examples, with labels)
<code>income.dev.csv</code>	dev set (1,000 examples, with labels)
<code>income.test.blind.csv</code>	(semi-blind) test set (1,000 examples, <u>without</u> labels)
<code>toy.csv</code>	toy data (only the first two fields)
<code>validate.py & random_output.py</code>	tools to validate your test result
<code>sample_submission.csv</code>	example submission with randomly generated (~50% positive) labels

The semi-blind test set does not contain the target labels (>50K or <=50K), which you will need to predict using your best model. Part of your grade is based on your prediction accuracy on test (see Part 5). You must use `validate.py` to validate your test result before submitting to Kaggle and Canvas.

You can also download it on a Linux/Mac command-line (e.g., on server):

```
$ wget http://eecs.oregonstate.edu/~huanlian/teaching/ML/2025fall/unit1/hw1/hw1-data.tgz
```

7. You should submit on Canvas separate files including `report.pdf`, `test.predicted.csv`, and all your code (either `.py` or `.ipynb`). **Do not forget the debrief section (in each HW).** `LATEX`'ing is recommended but not required.

1 Hands-on Exploration of the Income Dataset (3 pts) ([video 1](#) and [video 2](#))

Note: these questions in Part 1 can be done in (at least) two ways:

- **Linux/Mac terminals** (you can also use Jupyter Notebook to open a terminal or simulate a terminal in regular Python cells by ! such as `!ls`). This is our default method which is a useful tool for data scientists. If you're not familiar with terminals, please refer to the [Linux book](#) above and watch [video 1](#).
- **pandas**. If you have trouble with terminals, you can alternatively use the data processing module `pandas` which we will also use starting from Part 2. Please read the tutorial in Part 2 questions 1-2 and watch [video 2](#).

Questions (for each question, we provided hints for both methods; you can use either one but are recommended to try both in order to verify your answer and increase your understanding):

1. Take a look at the data. A training example (in this case, a negatively labeled one) looks like this:

`0,37,Private,Bachelors,Separated,Other-service,White,Male,70,England,<=50K`

which includes the ID column, the 9 input fields, and one output field (y):

`id, age, sector, education, marital-status, occupation, race, sex, hours-per-week, country-of-origin, target`

Q: What percentage of the training data has a positive label ($>50K$)? (This is known as **the positive %**). What about the dev set? Does it make sense given your knowledge of the average US per capita income? (0.5 pts)

Hint 1 (terminals): `$ grep -c ">50K" income.train.5k.csv` to count the positive people.

Note the \$ is the prompt, **not** part of the command (in jupyter notebook, use ! to run a command).

Hint 2 (pandas): `data[data["target"] == ">50K"]` to select the positive subset.

2. Q: What are the youngest and oldest ages in the training set? What are the least and most amounts of hours per week do people in the training set work? (0.5 pts)

Hint 1 (terminals): `$ cat income.train.5k.csv | tail -n +2 | sort -t, -nk2 | head -1`

Note: This command processes the CSV file as follows: `cat` lists the file contents (`cat` for “catalog”), `tail -n +2` skips the header by outputting from the second line onward, `sort -t, -nk2` sorts the lines numerically (-n) based on the second field (-k2) using comma as the separator (-t,), and `head -1` displays only the first line of the sorted output. The vertical bars (|) pipe each command’s output as input to the next.

Hint 2 (pandas): `max(data["age"])`.

3. The most important step in ML is **data preprocessing**, which maps **fields** (original data) to **features** (internal representation); see Parts 2–3 for details. There are two types of fields, *numerical* (`age` and `hours-per-week`), and *categorical* (everything else).¹ The default preprocessing method is to **binarize** all categorical fields, e.g., the **field race** becomes many **binary features** such as `race=White`, `race=Black`, etc. These features are binary because their values can only be 0 or 1, and for each example, in each field, there is one and only one positive feature (this is the so-called **“one-hot” representation**, widely used in ML and NLP).

Q: Why do we need to binarize all categorical fields? (0.5 pts)

Q: Why we do **not** want to binarize the two numerical fields, `age` and `hours`? (0.5 pts)

4. Q: How should we deal with the two numerical fields? Just as is, or normalize them (say, `age / 100`)? (0.5 pts)

Hint: The max Manhattan distance between two people on a categorial field is 2. If we simply normalize a numerical field (`age/100`), what’s the max distance on `age`? Are we treating all fields equally? See Part 3 Q2.

¹In principle, we could also convert `education` to a numerical feature, but we choose **not** to do it to keep it simple.

5. Q: How many features do you have in total (i.e., the dimensionality)? Hint: around 90. (Instead, if you binarize all fields, it would be around 230; see Part 2, Q3). **Note:** id and target are not features. (0.5 pts)

Hint 1 (terminals):

```
$ cat income.train.5k.csv | tail -n +2 | cut -f 3 -d ',' | sort | uniq -c
```

Here `cut` extracts specific columns of each line with `-f 3` for the 3rd column, and `-d ','` specifies the comma as the column separator. `uniq -c` displays unique rows with counts (e.g., 150 Federal-gov, 340 Local-gov, etc.). You can add `wc -l` to the pipeline to count the number of unique features (`wc` stands for “word count” and `-l` counts lines). You can also wrap this command with a for loop (`for i in `seq 2 10`; do ...; done`).

Hint 2 (pandas): use `pd.get_dummies()`; please refer to the tutorial in Part 2 question 2 for details.

2 Data Preprocessing and Feature Extraction I: Naive Binarization (3.5 pts)

In this section, we'll delve into a simple method of data preprocessing, *naive binarization*, and examine its implications and utility when applied to the Income dataset. Given the structure of our dataset, we have both numerical and categorical data. For the purpose of this exploration, we'll treat the numerical data (`age` and `hours-per-week`) equivalently to the categorical data. This means that `age=37` will be treated similarly to `sector=Private`.

1. **Pandas and Data Loading.** Before we proceed with feature extraction, let's understand how to load our dataset using the `pandas` library. The `read_csv` function facilitates this, and here we showcase loading from the toy dataset `toy.csv` (watch [video 2](#)):

```
import pandas as pd
data = pd.read_csv("toy.csv") # load the toy dataset
```

If you `print(data)`, it should output:

	<code>id</code>	<code>age</code>	<code>sector</code>
0	0	45	Federal-gov
1	1	33	Private
2	2	30	Local-gov
3	3	37	Private
4	4	35	Private
5	5	33	Private
6	6	40	Private
7	7	54	Federal-gov

(a) Pandas DataFrame Column Selection

i. Selecting a Single Column

There are two methods to select a single column:

- A. Method 1: You can specify the column inside a list within the DataFrame.

```
data_age1 = data[['age']]
```

- Result type: `pandas.core.frame.DataFrame`
- Shape: $(n, 1)$ where n is the number of rows
- 2-dimensional (table-like structure)

- B. Method 2: Including the column in the DataFrame directly.

```
data_age2 = data['age']
```

- Result type: `pandas.core.series.Series`
- Shape: $(n,)$ where n is the number of rows
- 1-dimensional

ii. Selecting Multiple Columns

To create a new DataFrame with a subset of columns, Pandas allows you to specify the necessary columns inside a list-like structure:

```
data_subset = data[['age', 'sector']]
```

Output of `print(data_subset):`

	age	sector
0	45	Federal-gov
1	33	Private
...		
6	40	Private
7	54	Federal-gov

- (b) **Masking** in pandas is useful for filtering data based on specific conditions. For instance, to view rows where `age` is 33 and `sector` is `Private`, use:

```
filtered_data = data[(data['age'] == 33) & (data['sector'] == 'Private')]
```

This returns a DataFrame containing only the rows that satisfy both conditions.

- 2. Binarization in Pandas.** Let's now introduce the concept of one-hot encoding with a practical example from our toy dataset. One-hot encoding is a technique to convert categorical data into a format that could be fed into machine learning algorithms. Essentially, for each unique value in a category, a binary feature is created. Using pandas, one-hot encoding can be done using `pandas.get_dummies()`. In the `columns` argument, which is used to specify which columns to binarize, we will only include the features `age` and `sector`, as `id` is not a feature.

```
encoded_data = pd.get_dummies(data, columns=["age", "sector"])
print(encoded_data)
```

This produces:

	id	age_19	age_30	age_33	age_35	age_37	age_40	age_45	age_47	age_54	sector_Federal-gov	sector_Local-gov	sector_Private
0	0	0	0	0	0	0	1	0	0	0	1	0	0
1	1	0	0	1	0	0	0	0	0	0	0	0	1
2	2	1	0	0	0	0	0	0	0	0	0	0	1
3	3	0	0	0	0	0	0	0	1	0	0	0	1
4	4	0	1	0	0	0	0	0	0	0	0	1	0
5	5	0	0	0	0	1	0	0	0	0	0	0	1
6	6	0	0	0	1	0	0	0	0	0	0	0	1
7	7	0	0	1	0	0	0	0	0	0	0	0	1
8	8	0	0	0	0	0	1	0	0	0	0	0	1
9	9	0	0	0	0	0	0	0	0	1	1	0	0

This output showcases the one-hot encoded version of our data. Each unique age and sector has been turned into a binary feature occupying its own column. In the resulting dataframe, if an individual's age is 45, then the `age_45` feature will be on (1), while all other age features will be off (0). Same for the `sector` features.

Notice how these binary features are named: the naming convention starts with the original field's name (e.g., `age` or `sector`), followed by an underscore (`_`) and then the unique value (e.g., `19` or `Federal-gov`). This naming rule is necessary: You might wonder why not just `19` instead of `age_19`? Well, there is also `hours_19`!

Question: Although `pandas.get_dummies()` is very handy for one-hot encoding, it's **absolutely impossible** to be used in machine learning. Why? (0.5 pts) (Hint: It's important to think about the entire pipeline. When working with training, dev and test sets, we need to ensure consistent representation across all of them.)

- 3. Binarization in Scikit-Learn.** A better tool is the `OneHotEncoder` class from the `sklearn.preprocessing` library. It ensures consistent representation across training, dev, and test sets because it can "fit" an encoder to the training set, which is then applied to all datasets. In other words, it's "stateful", as opposed to the "stateless" `pandas.get_dummies()`. It's slightly more complex than `pandas.get_dummies()` but is worth it:

- Unlike `pd.get_dummies()`, where you can select which columns to binarize, `OneHotEncoder` does not offer this option. Therefore, you must remove the `id` column from the DataFrame before applying one-hot encoding binarization.

```
data2 = data[['age', 'sector']]
```

- Import the necessary libraries:

```
from sklearn.preprocessing import OneHotEncoder
```

- Instantiate the encoder:

```
encoder = OneHotEncoder(sparse_output=False, handle_unknown='ignore')
```

Here `handle_unknown='ignore'` is important, otherwise when applied to the unseen dev/test data, it will throw an error on new features that it has not seen on the training set (e.g., a new country). See also Q2.

- Fit the encoder to the training data, and transform the training data:

```
encoder.fit(data2) # you only fit the encoder once (on training)
binary_data = encoder.transform(data2) # but use it to transform training, dev, and test set
```

In practice, you can combine these two lines by one function `encoder.fit_transform(data2)` but we chose to separate them for clarity. The output, `binary_data`, looks very similar to the binarized data from `pandas`:

```
[[0. 0. 0. 0. 0. 0. 1. 0. 0. 1. 0. 0.]
 [0. 0. 1. 0. 0. 0. 0. 0. 0. 0. 0. 1.]
 [1. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 1.]
 [0. 0. 0. 0. 0. 0. 0. 1. 0. 0. 0. 1.]
 [0. 1. 0. 0. 0. 0. 0. 0. 0. 0. 1. 0.]
 [0. 0. 0. 0. 1. 0. 0. 0. 0. 0. 0. 1.]
 [0. 0. 0. 1. 0. 0. 0. 0. 0. 0. 0. 1.]
 [0. 0. 1. 0. 0. 0. 0. 0. 0. 0. 0. 1.]
 [0. 0. 0. 0. 0. 1. 0. 0. 0. 0. 0. 1.]
 [0. 0. 0. 0. 0. 0. 1. 0. 0. 0. 0. 1.]]
```

You can check the name of each column by `encoder.get_feature_names_out()`, which gives us:

```
['age_19' 'age_30' 'age_33' 'age_35' 'age_37' 'age_40' 'age_45' 'age_47' 'age_54'
 'sector_Federal-gov' 'sector_Local-gov' 'sector_Private']
```

Question: After implementing the naive binarization to the real training set (NOT the toy set), what is the feature dimension? Does it match with the result from Part 1 Q5? (0.25 pts) Note: target is not a feature!

4. **Fit k -NN via Scikit-Learn.** With the dataset now binarized using the `OneHotEncoder`, an intriguing exploration is to employ the k -nearest neighbors (k -NN) algorithm on this transformed data. The k -NN algorithm works by classifying a data point based on how its neighbors are classified. The number of neighbors, denoted as k , is a parameter that can greatly affect the performance of the k -NN algorithm.

To begin with, you would utilize the `KNeighborsClassifier` from `sklearn.neighbors`. Given the binary representation of the dataset, you can follow these general steps:

- Prediction: Predict the labels for both the training set (to get training accuracy) and the dev set.
- Evaluation: Calculate and compare the accuracy scores for the predictions on the training and dev sets.

Questions:

- Evaluate k -NN on both the training and dev sets and report the error rates and predicted positive rates for k from 1 to 100 (odd numbers only, for tie-breaking), e.g., something like:

```

k=1      train_err xx.x% (+:xx.x%)  dev_err xx.x% (+:xx.x%)
k=3      ...
...
k=99     ...

```

Q: what's your best error rate on dev? Which k achieves this best error rate? (Hint: 1-NN dev error should be $\sim 23\%$ and the best dev error should be $\sim 16\%$). (1 pt)

- (b) Q: When $k = 1$, is training error 0%? Why or why not? Look at the training data to confirm your answer. (0.5 pts)
- (c) Q: What trends (train and dev error rates and positive ratios, and running speed) do you observe with increasing k ? (0.75 pts)
- (d) Q: What does $k = \infty$ actually do? Is it extreme overfitting or underfitting? What about $k = 1$? (0.5 pts)
- (e) Using your best model (in terms of dev error rate), predict the semi-blind test data, and submit it to Kaggle (follow instructions from Part 5). What are your error rate and ranking on the public leaderboard? Take a screenshot. Hint: your public error rate should be $\sim 18.5\%$. (extra 0.5 pts)

3 Data Preprocessing and Feature Extraction II: Smart Binarization (3 pts)

In the previous section, we discussed the naive binarization of the Income dataset. While this method can be effective for certain machine learning models, there are more refined data preprocessing techniques that can yield better results, especially when dealing with a mix of numerical and categorical data. This section delves into smarter binarization methods, particularly focusing on differential handling of numerical and categorical features and scaling the numerical data for better algorithmic performance.

- For machine algorithms sensitive to distance metrics such as k -NN, it's pivotal to preprocess data more intelligently. In our case, we want to treat numerical (`age` and `hours-per-week`) and categorical fields differently.

For categorical data, we continue to use `OneHotEncoder` for binarization. However, for numerical data, instead of binarizing them directly, we let them remain in their original format. Here's how you'd typically proceed:

- Differentiate between numerical and categorical fields in your dataset (i.e., identify their column names).
- Apply `OneHotEncoder` only to the categorical fields (in our case, all but `age` and `hours-per-week`).
- Combine the processed categorical data with the numerical data to form the modified dataset.

Here, we provide an example for the toy dataset:

- Import the necessary libraries:

```

from sklearn.compose import ColumnTransformer
from sklearn.preprocessing import OneHotEncoder

```

- Instantiate the preprocessing methods for numerical and categorical data:

```

num_processor = 'passthrough' # i.e., no transformation
cat_processor = OneHotEncoder(sparse_output=False, handle_unknown='ignore')

```

- Fit the processor to the dataset and transform:

```

processor = ColumnTransformer([
    ('num', num_processor, ['age']),
    ('cat', cat_processor, ['sector'])
])
processor.fit(data)
processed_data = processor.transform(data)

```

Note: The `preprocessor.transform(data)` will only return the columns which are included in `ColumnTransformer`. As `id` is not included in it, we can use the original `data` which contains the `id` column; so this is easier than `OneHotEncoder`. The `processed_data` will contain only the transformed columns as:

```
[[45.  1.  0.  0.]
 [33.  0.  0.  1.]
 [19.  0.  0.  1.]
 [47.  0.  0.  1.]
 [30.  0.  1.  0.]
 [37.  0.  0.  1.]
 [35.  0.  0.  1.]
 [33.  0.  0.  1.]
 [40.  0.  0.  1.]
 [54.  1.  0.  0.]]
```

You can also check the name of each feature by `preprocessor.get_feature_names_out()`, which gives us:

```
['num_age' 'cat_sector_Federal-gov' 'cat_sector_Local-gov' 'cat_sector_Private']
```

After this preprocessing step, make your new number of features is around 90, as in Part 1 Q5.

Question: Re-execute all experiments with varying values of k (Part 2, Q 4a) and report the new results. Do you notice any performance improvements compared to the initial results? If so, why? If not, why do you think that is? (1.5 pts) (Hint: 1-NN dev error should be $\sim 27\%$ and the best dev error should be $\sim 21\%$)

- Simply keeping numerical data unchanged may not always yield optimal results. Algorithms like k -NN, which rely on distance metrics, can be sensitive to the scale of the features. A variable that spans a large range can unduly influence the distance computation. For instance, `age` and `hours-per-week` typically have much larger ranges than other binary features, which could result in the two of them dominating the distance calculations.

To tackle this, we can perform rescaling on the numerical fields, ensuring they lie within a similar range, e.g., bounding each field's max Manhattan distance by 2 (see Part 1 Q4). `MinMaxScaler` from `sklearn.preprocessing` is a tool designed for this task, scaling features to lie between a given range, often between zero and one.

Here's an improved version of the previous implementation:

- Import the necessary libraries:

```
from sklearn.compose import ColumnTransformer
from sklearn.preprocessing import MinMaxScaler, OneHotEncoder
```

- Instantiate the processing methods for numerical and categorical data:

```
num_processor = MinMaxScaler(feature_range=(0, 2))
cat_processor = OneHotEncoder(sparse_output=False, handle_unknown='ignore')
```

- Fit the processor to the dataset and transform:

```
preprocessor = ColumnTransformer([
    ('num', num_processor, ['age']),
    ('cat', cat_processor, ['sector'])
])
preprocessor.fit(data)
processed_data = preprocessor.transform(data)
```

This new version will produce more reasonable features for numerical data:

```
[[1.48571429  1.          0.          0.          ]
 [0.8           0.          0.          1.          ]
 [0.           0.          0.          1.          ]]
```

```
[1.6      0.      0.      1.      ]
[0.62857143 0.      1.      0.      ]
[1.02857143 0.      0.      1.      ]
[0.91428571 0.      0.      1.      ]
[0.8      0.      0.      1.      ]
[1.2      0.      0.      1.      ]
[2.      1.      0.      0.      ]]
```

Questions: Again, rerun all experiments with varying values of k and report the results (Part 2, Question 4a). Do you notice any performance improvements? If so, why? If not, why do you think that is? (1.5 pts) (Hint: 1-NN dev error should be $\sim 24\%$ and the best dev error should be $\sim 14\text{--}15\%$)

- Again, using your new best model (in terms of dev error rate), predict the semi-blind test data, and submit it to Kaggle. What are your new error rate and ranking on the public leaderboard? Take a screenshot. Hint: your public error rate should improve to $\sim 17\%$. (extra 0.5 pts)

4 Implement your own k -Nearest Neighbor Classifiers (4.5 pts)

While leveraging libraries like `sklearn` can be incredibly convenient for employing algorithms like k -NN, there's significant educational value in implementing the algorithm from scratch by yourself. By doing so, you can gain a deeper understanding of the intricacies and operations that occur under the hood.

Note: you should start with the preprocessed data from `sklearn` (smart binarization + scaling), and only implement the k -NN classifier; i.e., no need to implement the binarization in Python yourself, which is non-trivial.

Note: you can use the Matlab style “broadcasting” notations in numpy (such as matrix - vector) to calculate many distances in one shot. For example, if A is an $n \times m$ matrix (n rows, m columns, where n is the number of people and m is the number of features), and p is an m -dimensional vector (1 row, m columns) representing the query person, then $A - p$ returns the difference vectors from each person in A to the query person p , from which you can compute the distances:

```
>>> A = np.array([[1,2], [2,3], [4,5]]); p = np.array([3,2])
>>> A - p
array([[-2,  0],     [-1,  1],     [ 1,  3]])
>>> np.linalg.norm(A-p, axis=1)
array([2.          , 1.41421356, 3.16227766])
```

This is Euclidean distance (what does `axis=1` mean?). You need to figure out Manhattan distance yourself.

- Distance Verification (sanity check).** Once you have prepared the dataset, use the `sklearn` k -NN predictor (with rescaled smart binarization) to validate your results for the first person in the dev dataset. Particularly, identify the three closest individuals from the training set. The first person in the dev set is:

```
$ head -1 income.dev.csv
5000,45,Federal-gov,Bachelors,Married-civ-spouse,Adm-clerical,White,Male,45,United-States,<=50K
```

Its binarized features should look like:

```
age sector_Federal-gov ... occupation_Adm-clerical ... hours ... country_United-States ...
0.767123           1.0 ...                   1.0 ... 0.897959 ...           1.0 ...
```

Using the `kneighbors` method from `KNeighborsClassifier`, according to the Euclidean distance, you can find the following three closest individuals from the training dataset (with 0-based row index):

```
(4872) age sector_Federal-gov ... occupation_Adm-clerical ... hours ... country_United-States ...
0.438356           1.0 ...                   1.0 ... 0.836735 ...           1.0 ...
(4787) age sector_Federal-gov ... occupation_Adm-clerical ... hours ... country_United-States ...
0.821918           1.0 ...                   1.0 ... 0.897959 ...           0.0 ...
(2591) age sector_Federal-gov ... occupation_Adm-clerical ... hours ... country_United-States ...
0.849315           1.0 ...                   0.0 ... 0.877551 ...           1.0 ...
```

Their Euclidean distances to the query person are approximately 0.334, 1.415, and 1.417 (make sure your `sklearn` results match these). The original rows in the `train` dataset:

```
$ grep '^4872' income.train.5k.csv
4872,33,Federal-gov,Bachelors,Married-civ-spouse,Adm-clerical,White,Male,42,United-States,>50K
$ grep '^4787' income.train.5k.csv
4787,47,Federal-gov,Bachelors,Married-civ-spouse,Adm-clerical,White,Male,45,Germany,>50K
$ grep '^2591' income.train.5k.csv
2591,48,Federal-gov,Bachelors,Married-civ-spouse,Prof-specialty,White,Male,44,United-States,>50K
```

In the `grep` command, the "`^`" in "`^4872`" specifies that "4872" should be found at the start of the line. Here is the detailed computation of distances for your convenience:

```
sqrt((0.767123-0.438356)**2 + (0.897959-0.836735)**2) = 0.334          # only differ in age and hours
sqrt((0.767123-0.821918)**2 + 1**2 + 1**2) = 1.415                      # only differ in age and country
sqrt((0.767123-0.849315)**2 + 1**2 + 1**2 + (0.897959-0.877551)**2) = 1.417 # age, occupation, and hours
```

Also notice that in this example, the 3-NN prediction is wrong, as the top-3 closest examples are all `>50K`.

Finally, remember that you don't really need to sort the distances in order to get the top- k closest examples.

Note: If you couldn't get the same top-3 people listed above, it is possible that you included the target label in your feature map.

Question: Before implementing your k -NN classifier, try to verify the distances from your implementation with those from the `sklearn` implementation. What are the (Euclidean and Manhattan) distances between the query person above (the first in dev set) and the top-3 people listed above? Report results from both `sklearn` and your own implementation. (0.5 pts)

2. Implement your own k -NN classifier (with the default Euclidean distance).

- (a) Q: Is there any work in training after the feature map (i.e., after all fields become features)? (0.25 pts)
- (b) Q: What's the time complexity of k -NN to test one example (dimensionality d , size of training set $|D|$)? (0.75 pt) If you do not have a CS background, please state it here.
- (c) Q: Do you really need to sort the distances first and then choose the top k ? Hint: there is a faster way to choose top k without sorting. (0.5 pts)
- (d) Q: What numpy tricks did you use to speed up your program so that it can be fast enough to print the training error? Hint: (i) broadcasting (such as matrix - vector); (ii) `np.linalg.norm(..., axis=1)`; (iii) `np.argsort()` or `np.argpartition()`; (iv) slicing. The main idea is to do as much computation in the vector-matrix format as possible (i.e., the Matlab philosophy), and as little in Python as possible. (1 pt)
- (e) Q: How many seconds does it take to print the training and dev errors for $k = 99$ on ENGR servers? Hint: use `$ time python ...` and report the user time instead of real time. (Mine was 14 seconds). (0.5 pts)

3. Redo the evaluation using Manhattan distance (for $k = 1 \sim 99$). Better or worse? (1 pt)

5 Deployment (3.5 pts)

Let's try more k 's and take your best model (according to dev error rate) and run it on the semi-blind test data, and produce `income.test.predicted.csv`, which has the same format as the training and dev files.

Q1: At which k and with which distance did you achieve the best dev results? (0.25 pts)

Q2: What's your best dev error rates and the corresponding positive ratios? (0.25 pts)

Q3: What's the positive ratio on test? (0.25 pts)

Part of your grade will depend on the accuracy of `income.test.predicted.csv` (2.75 pts).

IMPORTANT: You should use our `validate.py` (from `hw1-data.tgz`) to verify your `income.test.predicted.csv`; it will catch many common problems such as formatting issues and overly positive or overly negative results:

```
$ cat income.test.predicted.csv | python3 validate.py
```

We also provided a `random_output.py` which generates random predictions (by default ~50% positive; see `sample_submission.csv` as an example) and it will pass the formatting check, but fail on the positive ratio:

```
$ cat income.test.blind.csv | python3 random_output.py | python3 validate.py
```

which might output:

```
Your file passed the formatting test! :)
```

```
Your positive rate is 49.6%.
```

```
ERROR: Your positive rate seems too high (should be similar to train and dev).
```

```
PLEASE DOUBLE CHECK YOUR BINARIZATION AND kNN CODE.
```

However you randomly generate ~25% positive labels and it will pass both checks:

```
$ cat income.test.blind.csv | python3 random_output.py 0.25 | python3 validate.py
```

Now submit your `income.test.predicted.csv` to the Kaggle competition, and get it scored. Check the leaderboard to see your rank.

Notice that:

- The public leaderboard is based on the first 500 examples in the test data, and there is a private leaderboard for the other 500 examples. Your grade will be based on both the **public** and **private** leaderboard scores. This is to reward hardworking while preventing overfitting. The private leaderboard includes all your submissions, not just your best one on the public board. It is possible that your non-best entry on the public board becomes your best on the private one.
- You can only submit up to **10 times per day**. This is also to prevent overfitting.
- You have to use your **official OSU email** to create an account for this competition (and all other HWs in this course), and your “**team name**” (after joining the competition) must be the **last 5 digits of your OSU ID**. We will announce the top performers (in both public and private leaderboards) after HW1 is due.
- You are only allowed to use **one** account for this competition. This is to prevent overfitting and cheating.
- You can only use *k*-NN for this HW. Do **not** use any other machine learning algorithms (to ensure fairness).

Q4: What's your best rank on the public leaderboard? Take a screenshot. How many submissions did you use?

6 Observations (2.5 pts)

1. Q: Summarize the major drawbacks of *k*-NN that you observed by doing this HW. There are a lot! (1 pt) Hint: The most obvious one is: Why are all fields treated equally? Which field should be more important?
2. Q: Do you observe in this HW that best-performing models tend to exaggerate the existing bias in the training data? Is it due to overfitting or underfitting? Is this a potentially social issue? (1.5 pts) Hint: for example, compare the true positive % vs. your predicted positive % on the dev set. What about the positive % given gender? (e.g., for all females in the dev set, compare the true positive % and your predicted one.) Or race?

Note: All machine learning algorithms tend to exaggerate existing biases in data. Why? Because learning involves compression (see https://en.wikipedia.org/wiki/Data_compression#Machine_learning). This is like when you really understand a textbook, you've effectively compressed it to a few bullet points. Machine learning, like human learning, has to compress the data to a smaller representation, otherwise it's just “memorizing the textbook”, which is useless at testing time (overfitting). To do this compression, machine learning has to sacrifice some details. In *k*-NN, even though there is no training per se, the use of a larger *k* effectively draws a smoother boundary between positive and negative regions (thus smaller representation of the concept classes). By contrast, *k* = 1 draws the most zigzagged boundary (see slides). But this is a serious social issue because ML is used everywhere.

7 Extra Credit Question (extra 2 pts)

For an additional challenge, visualize the dev error rates of all three versions (naive, smart, and smart+scaling, all with Euclidean distance) against k . Plot these three error rate curves on a single graph to allow for easy comparison. Use different colors or styles for each curve and be sure to include a legend to distinguish between them.

Debriefing (required in your report)

1. Approximately how many hours did you spend on this assignment?
2. Would you rate it as easy, moderate, or difficult?
3. Did you work on it mostly alone, or mostly with other people?
4. How deeply do you feel you understand the material it covers (0%–100%)?
5. Any other comments?